

August 7, 2018

1 DMG2 Assignment : Problem 5

k-Nearest Neighbours Classifier, Parzen Window Classifier

```
In [1]: import pandas as pd
import numpy as np
import os
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.neighbors import KNeighborsClassifier
from sklearn import preprocessing
from sklearn.neighbors import KernelDensity

sns.set_style('whitegrid')

In [2]: DATA_DIR = '/home/jishnu/Documents/ISB/Term3/dmg2/assignments/hw_assignment1/dmg2/data'

In [3]: train = pd.DataFrame(columns=['V{}'.format(i) for i in range(1,785)] + ['label'])
test = pd.DataFrame(columns=['V{}'.format(i) for i in range(1,785)] + ['label'])
for num in range(10):
    # Consolidating training data
    temp_train = pd.read_csv(os.path.join(DATA_DIR, 'train{0}.csv'.format(num)), usecols=
    temp_train['label'] = num
    train = train.append(temp_train, ignore_index=True)
    # Consolidating test data
    temp_test = pd.read_csv(os.path.join(DATA_DIR, 'test{0}.csv'.format(num)), usecols=
    temp_test['label'] = num
    test = test.append(temp_test, ignore_index=True)

In [4]: train.shape

Out[4]: (36470, 785)

In [5]: test.shape
```

```
Out[5]: (24190, 785)
```

```
In [6]: train[train.isnull().any(axis=1)].groupby(by='label')['label'].value_counts()
```

```
Out[6]: label  label
         4         46
         5        299
         6         1
         8         42
        Name: label, dtype: int64
```

```
In [7]: test[test.isnull().any(axis=1)].groupby(by='label')['label'].value_counts()
```

```
Out[7]: label  label
         4         35
         5        203
         6         4
         8         30
        Name: label, dtype: int64
```

There are missing values in both the training and test data. Shown above is the count of rows with missing values, and the associated labels.

```
In [8]: train.groupby(by='label')['label'].value_counts()
```

```
Out[8]: label  label
         0        3567
         1        4034
         2        3582
         3        3677
         4        3567
         5        3567
         6        3567
         7        3763
         8        3567
         9        3579
        Name: label, dtype: int64
```

```
In [9]: test.groupby(by='label')['label'].value_counts()
```

```
Out[9]: label  label
         0        2356
         1        2708
         2        2376
         3        2454
         4        2356
         5        2356
         6        2356
         7        2502
         8        2356
         9        2370
        Name: label, dtype: int64
```

Considering the number of complete data for each label, we can safely remove the rows with missing values for our analysis.

```
In [10]: train = train.dropna()
        test = test.dropna()
```

```
In [11]: train.isnull().values.any()
```

```
Out[11]: False
```

```
In [12]: test.isnull().values.any()
```

```
Out[12]: False
```

There are no missing values in the training and test data now

```
In [13]: X_train = train.iloc[:, :784]
        Y_train = train.iloc[:, 784]
```

```
X_test = test.iloc[:, :784]
```

```
Y_test = test.iloc[:, 784]
```

```
In [14]: # Standardizing feature values
        X_train = StandardScaler().fit_transform(X_train)
        X_test = StandardScaler().fit_transform(X_test)
```

1.1 Applying PCA

```
In [15]: # Applying PCA
        pc = PCA(n_components=9).fit_transform(X_train)
        d1_train = pd.DataFrame(data=pc, columns=['pc{0}'.format(i) for i in range(1,10)])
        d1_train['label'] = Y_train.values
        d1_train.head(5)

        pc = PCA(n_components=9).fit_transform(X_test)
        d1_test = pd.DataFrame(data=pc, columns=['pc{0}'.format(i) for i in range(1,10)])
        d1_test['label'] = Y_test.values
        d1_test.head(5)
```

```
Out[15]:
```

	pc1	pc2	pc3	pc4	pc5	pc6	pc7	\
0	1.751833	-6.389755	-2.021087	-2.694718	-6.427682	1.025097	-0.554033	
1	5.884343	-7.690717	-2.390982	0.260438	-4.924087	-0.382575	0.298304	
2	16.381093	5.663541	-1.865832	-3.368391	4.353517	-3.098186	-5.629483	
3	12.814322	-7.638621	-4.434501	-7.660043	-0.550385	-1.785665	-0.076132	
4	11.123304	7.252218	5.007200	-0.333093	15.020012	-0.071782	-0.972011	

	pc8	pc9	label
0	5.296349	3.807623	0
1	6.510170	3.098220	0
2	-3.894353	2.471996	0
3	1.813952	0.309337	0
4	-2.298600	0.127185	0

1.2 Applying Fisher LDA

```
In [16]: fisher = LinearDiscriminantAnalysis(n_components=9).fit_transform(X_train,Y_train.astype('float32'))
d2_train = pd.DataFrame(data=fisher,columns=['f{0}'.format(i) for i in range(1,10)])
d2_train['label'] = Y_train.values
d2_train.head(5)

fisher = LinearDiscriminantAnalysis(n_components=9).fit_transform(X_test,Y_test.astype('float32'))
d2_test = pd.DataFrame(data=fisher,columns=['f{0}'.format(i) for i in range(1,10)])
d2_test['label'] = Y_test.values
d2_test.head(5)

/home/jishnu/anaconda3/lib/python3.6/site-packages/sklearn/discriminant_analysis.py:388: UserWarning:
  warnings.warn("Variables are collinear.")
/home/jishnu/anaconda3/lib/python3.6/site-packages/sklearn/discriminant_analysis.py:442: UserWarning:
  UserWarning)
/home/jishnu/anaconda3/lib/python3.6/site-packages/sklearn/discriminant_analysis.py:388: UserWarning:
  warnings.warn("Variables are collinear.")
```

```
Out [16]:
```

	f1	f2	f3	f4	f5	f6	f7	\
0	-2.833687	-1.079133	-0.845838	-0.764987	-0.484096	0.313129	-0.955681	
1	-3.849851	-3.435085	-1.935596	0.478846	-2.420500	-0.475621	0.572627	
2	-3.350865	-4.500847	-3.900589	-1.086583	-2.654538	-1.739049	-1.565408	
3	-3.283216	-2.503597	-3.138834	-0.103059	-1.582149	-0.295842	1.330582	
4	-2.231976	-1.552025	-3.889057	-0.017441	-0.577388	-2.299944	-0.684894	

	f8	f9	label
0	-0.997196	1.076134	0
1	-0.989022	0.779083	0
2	-0.213139	-0.024020	0
3	-3.239352	1.473339	0
4	0.945457	0.404140	0

1.3 k-Nearest Neighbors Classification

```
In [17]: d1_train_X = d1_train.iloc[:, :9]
d1_train_Y = d1_train.iloc[:, 9].astype('int')

d1_test_X = d1_test.iloc[:, :9]
d1_test_Y = d1_test.iloc[:, 9].astype('int')

d2_train_X = d2_train.iloc[:, :9]
d2_train_Y = d2_train.iloc[:, 9].astype('int')

d2_test_X = d2_test.iloc[:, :9]
d2_test_Y = d2_test.iloc[:, 9].astype('int')

In [18]: d1_knn = pd.DataFrame(columns=['k', 'acc_type', 'acc'])
d2_knn = pd.DataFrame(columns=['k', 'acc_type', 'acc'])
```

```
In [19]: for k in range(1,18,2):
          knn1 = KNeighborsClassifier(n_neighbors=k).fit(d1_train_X,d1_train_Y)
          knn2 = KNeighborsClassifier(n_neighbors=k).fit(d2_train_X,d2_train_Y)
          d1_knn.append({'k' : k, 'acc_type' : 'training', 'acc' : np.round(knn1.score,3)})
          d1_knn.append({'k' : k, 'acc_type' : 'test', 'acc' : np.round(knn1.score,3)})
          d2_knn.append({'k' : k, 'acc_type' : 'training', 'acc' : np.round(knn2.score,3)})
          d2_knn.append({'k' : k, 'acc_type' : 'test', 'acc' : np.round(knn2.score,3)})
```

```
In [20]: d1_knn.head()
```

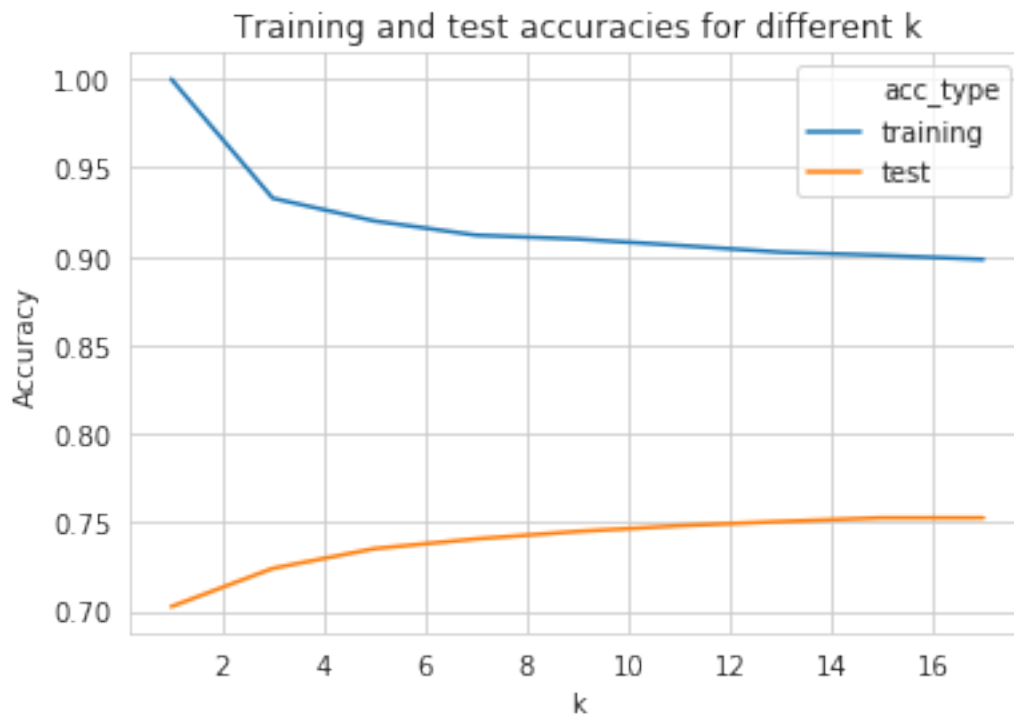
```
Out[20]:
```

	k	acc_type	acc
0	1	training	1.0000
1	1	test	0.7025
2	3	training	0.9328
3	3	test	0.7241
4	5	training	0.9200

1.3.1 Plotting training and test accuracy for kNN Classification

D1 Dataset

```
In [21]: sns.lineplot(x='k',y='acc',hue='acc_type',data=d1_knn,ci=0,)
          plt.title('Training and test accuracies for different k')
          plt.xlabel('k')
          plt.ylabel('Accuracy')
          plt.show();
```



D2 Dataset

```
In [22]: sns.lineplot(x='k',y='acc',hue='acc_type',data=d2_knn,ci=0,)
plt.title('Training and test accuracies for different k')
plt.xlabel('k')
plt.ylabel('Accuracy')
plt.show();
```



The optimal k for both datasets is 8 when considering the test accuracies.

1.4 Parzen-Window Classification

For each data point in test set, find the kernel function of the form $\exp(-(X_i - X_t)^2 / 2\sigma^2) / \sigma$

The distance function used in euclidean, and the sum of kernel function value is found for all training data for each class to come up with the score of the class.

This score is converted to a probability to find the predicted class with maximum probability

Sampling values from each class for better performance

```
In [23]: d1_train.groupby('label', group_keys=False).count()
```

```
Out[23]:
```

	pc1	pc2	pc3	pc4	pc5	pc6	pc7	pc8	pc9
label									
0	3567	3567	3567	3567	3567	3567	3567	3567	3567

1	4034	4034	4034	4034	4034	4034	4034	4034	4034
2	3582	3582	3582	3582	3582	3582	3582	3582	3582
3	3677	3677	3677	3677	3677	3677	3677	3677	3677
4	3521	3521	3521	3521	3521	3521	3521	3521	3521
5	3268	3268	3268	3268	3268	3268	3268	3268	3268
6	3566	3566	3566	3566	3566	3566	3566	3566	3566
7	3763	3763	3763	3763	3763	3763	3763	3763	3763
8	3525	3525	3525	3525	3525	3525	3525	3525	3525
9	3579	3579	3579	3579	3579	3579	3579	3579	3579

```
In [24]: d1_test.groupby('label', group_keys=False).count()
```

```
Out[24]:
```

	pc1	pc2	pc3	pc4	pc5	pc6	pc7	pc8	pc9
label									
0	2356	2356	2356	2356	2356	2356	2356	2356	2356
1	2708	2708	2708	2708	2708	2708	2708	2708	2708
2	2376	2376	2376	2376	2376	2376	2376	2376	2376
3	2454	2454	2454	2454	2454	2454	2454	2454	2454
4	2321	2321	2321	2321	2321	2321	2321	2321	2321
5	2153	2153	2153	2153	2153	2153	2153	2153	2153
6	2352	2352	2352	2352	2352	2352	2352	2352	2352
7	2502	2502	2502	2502	2502	2502	2502	2502	2502
8	2326	2326	2326	2326	2326	2326	2326	2326	2326
9	2370	2370	2370	2370	2370	2370	2370	2370	2370

There are around 3000 - 4000 data points for each class in train and 2000 - 3000 data points for each class in test. Let's sample 10 data points from each class

```
In [25]: d1_train = d1_train.groupby('label', group_keys=False).apply(lambda x: x.sample(min(10, len(x))))
d1_test = d1_test.groupby('label', group_keys=False).apply(lambda x: x.sample(min(10, len(x))))

d2_train = d2_train.groupby('label', group_keys=False).apply(lambda x: x.sample(min(10, len(x))))
d2_test = d2_test.groupby('label', group_keys=False).apply(lambda x: x.sample(min(10, len(x))))

d1_train_X = d1_train.iloc[:, :9]
d1_train_Y = d1_train.iloc[:, 9].astype('int')

d1_test_X = d1_test.iloc[:, :9]
d1_test_Y = d1_test.iloc[:, 9].astype('int')

d2_train_X = d2_train.iloc[:, :9]
d2_train_Y = d2_train.iloc[:, 9].astype('int')

d2_test_X = d2_test.iloc[:, :9]
d2_test_Y = d2_test.iloc[:, 9].astype('int')

In [26]: def d1_train_ker(row_x, row_y, sigma):
    dist = np.linalg.norm(row_x - row_y)
    kernel_fn = np.exp(-dist**2/(2*sigma**2)) / sigma
```